EE141-Spring 2010 Digital Integrated Circuits

Lecture 13
Complex CMOS - Revisited

Administrativia

- No more re-grading of Midterm after today
- Hw 5 due on Friday.
- Hw 6 posted early next week. You get TWO weeks for this one.
- Project phase 1 introduced today
  - Actual launch is on Friday
- Out of town next week
  - We lecture offered by Stanley
  - Fr lecture cancelled – Make-up on Tu March 16 at 3:30pm

Class Material

- Last lecture
  - Inverter delay
- Today's lecture
  - Inverter energy
  - Project launch
  - Optimizing complex CMOS
- Reading (Ch 5, 6)

Propagation Delay

Transient Response

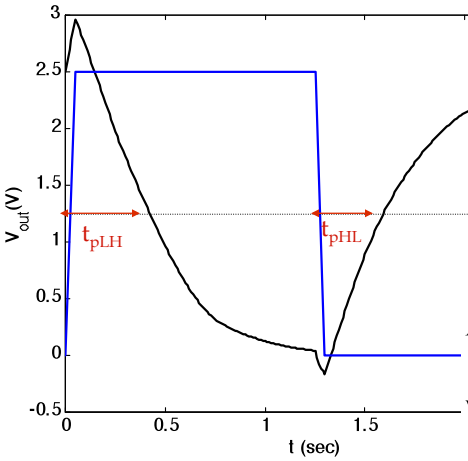

$t_{pHL} = 0.69\ C_L\ R_{eqn}$
$t_{pLH} = 0.69\ C_L\ R_{eqp}$

$$R_{eq} = \frac{1}{2\ln(2)}\frac{1}{(1+\lambda V_{DD})}\frac{V_{DD}}{I_{DSAT}}$$

$$\text{with } I_{DSAT} = k'\frac{W}{L}\left(V_{DD} - V_T - \frac{V_{DS,eff}}{2}\right)V_{DS,eff}$$

Step Inputs?

- Derived RC model assuming input was a step
  - But input is not a step
  - Transistor turns on gradually
- Let's look at gate switching more carefully
  - Use our models to understand the effect of input slope

Input Slope Dependence

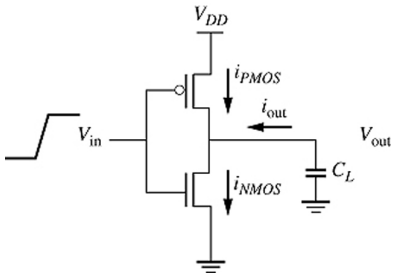

$$I_{out} = C_L \frac{dV_{out}}{dt} = I_{NMOS} - I_{PMOS}$$

- One way to analyze slope effect
  - Plug non-linear IV into diff. equation and solve…
- Simpler, approximate solution:
  - Use V_T* model

Slope Analysis

- For falling edge at output:
  - For reasonable inputs, can ignore I_{PMOS}
  - Either V_{ds} is very small, or V_{gs} is very small
- So, output current ramp starts when $V_{in}=V_T$*
  - Could evaluate the integral
  - Learn more by using an intuitive, graphical approach

Result Summary

- For reasonable input slopes:

$$t_{p,ramp} = t_{p,step} + \frac{V_T^*}{V_{DD}} \cdot t_{p,in}$$

Model vs. Spice Data

- For reasonable input slope
  - Model matches Spice very well
- Model breaks with very large t_r
  - Input looks "DC" – traces out VTC
  - Have other problems here anyways
    - Short-circuit current

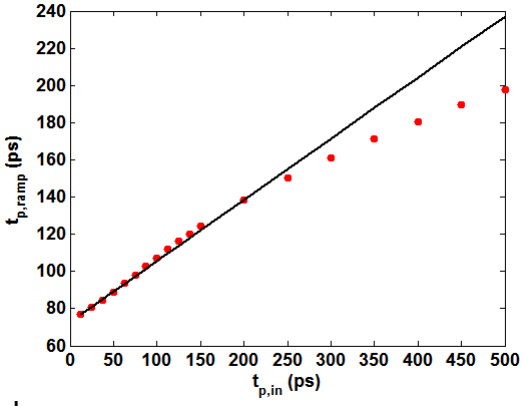

CMOS Power Dissipation

Where Does Power Go in CMOS?

- Switching power
  - Charging/discharging capacitors
- Leakage power
  - Transistors are imperfect switches
- Short-circuit power
  - Both pull-up and pull-down on during transition
- Static currents
  - Biasing currents, in e.g. analog, memory

Dynamic Power Consumption

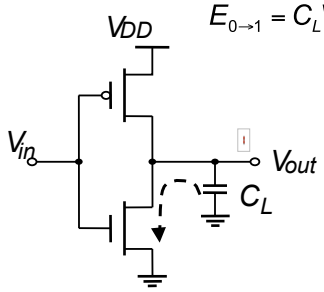

$$E_{0\to1} = C_L V_{DD}^2$$

$$E_R = \frac{1}{2} C_L V_{DD}^2$$

$$E_C = \frac{1}{2} C_L V_{DD}^2$$

- One half of the energy from the supply is consumed in the pull-up network, one half is stored on C_L
- Energy from C_L is dumped during the 1→0 transition

Dynamic Power Consumption

Power = Energy/transition • Transition rate

$$= C_L V_{DD}^2 \bullet f_{0\to1}$$

$$= C_L V_{DD}^2 \bullet f \bullet \alpha_{0\to1}$$

$$= C_{switched} V_{DD}^2 \bullet f$$

- Power dissipation is data dependent – depends on the switching probability
- Switched capacitance $C_{switched} = C_L \bullet \alpha_{0\to1}$

Transition Activity and Power

- Energy consumed in N cycles, E_N:

$$E_N = C_L \bullet V_{DD}^2 \bullet n_{0\to1}$$

$n_{0\to1}$ – number of 0→1 transitions in N cycles

$$P_{avg} = \lim_{N\to\infty} \frac{E_N}{N} \cdot f = \left(\lim_{N\to\infty} \frac{n_{0\to1}}{N} \right) \cdot C_L \cdot V_{DD}^2 \cdot f$$

$$\alpha_{0\to1} = \lim_{N\to\infty} \frac{n_{0\to1}}{N}$$

$$\boxed{P_{avg} = \alpha_{0\to1} \cdot C_L \cdot V_{DD}^2 \cdot f}$$

Short Circuit Current

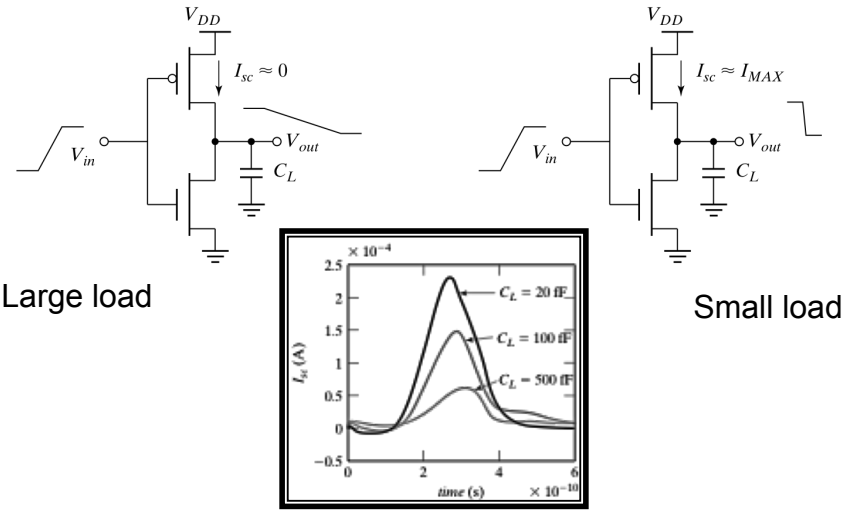

Large load

Small load

- Short circuit current usually well controlled

Diode Leakage

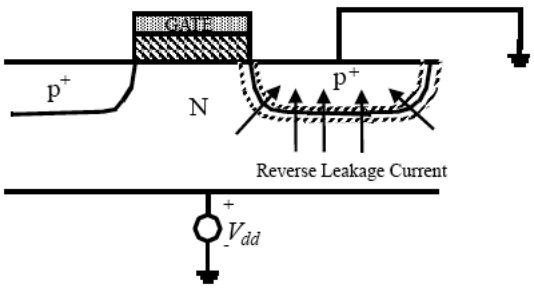

$$I_{DL} = J_S \times A$$

JS = 10-100 pA/μm2 at 25 deg C for 0.25μm CMOS
JS doubles for every 9 deg C!
Much smaller than transistor leakage in deep submicron

Transistor Leakage

- Transistors that are supposed to be off - leak

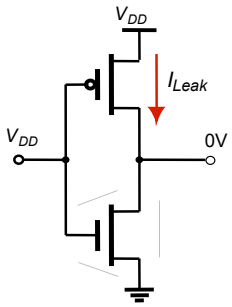

Input at V_{DD}

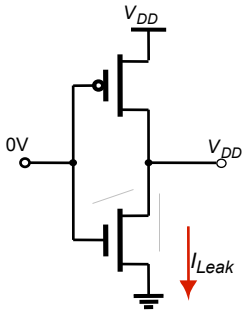

Input at 0

Transistor Leakage

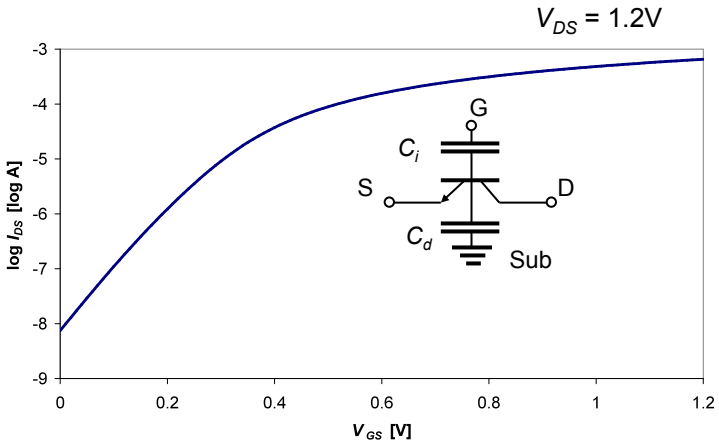

Drain leakage current is exponential with V_{GS}-V_T

Sub-Threshold Conduction

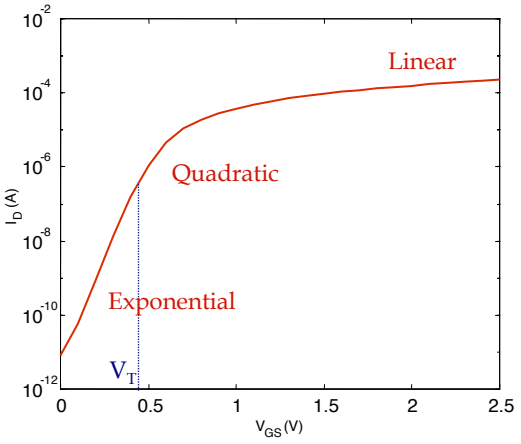

Subthreshold Slope:

$$I_D \sim I_0 e^{\frac{q(V_{GS}-V_T)}{nkT}}, \quad n = 1 + \frac{C_D}{C_{ox}}$$

S is ΔV_{GS} for I_{D2}/I_{D1} =10

$$S = n\left(\frac{kT}{q}\right)\ln(10)$$

Typical values for S:
60 .. 100 mV/decade

Transistor Leakage

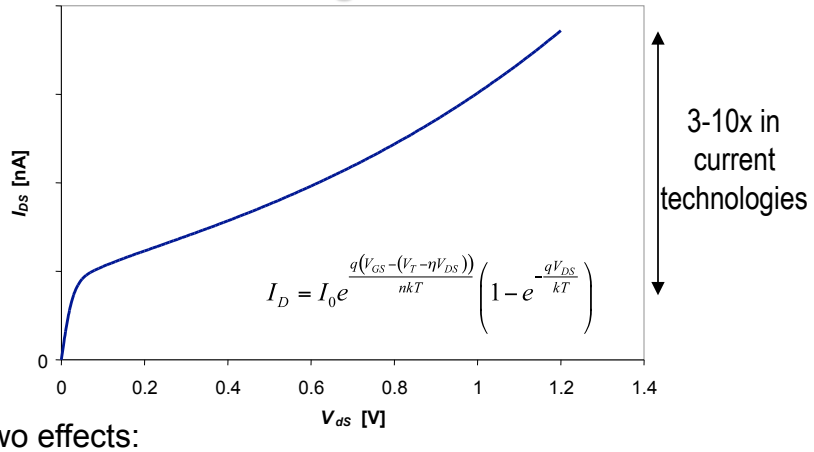

Two effects:
- diffusion current (like a bipolar transistor)
- exponential increase with V_{DS} (η: DIBL)

Threshold Variations

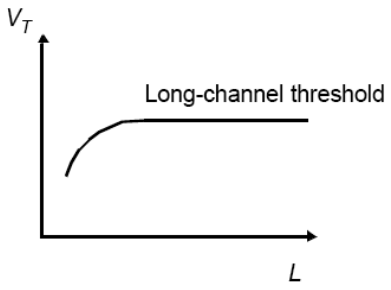

Threshold as a function of the length (for low V_{DS})

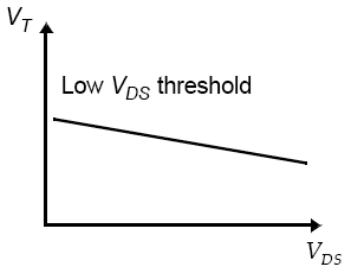

Drain-induced barrier lowering (DIBL) (for short L)

Trading off energy and delay

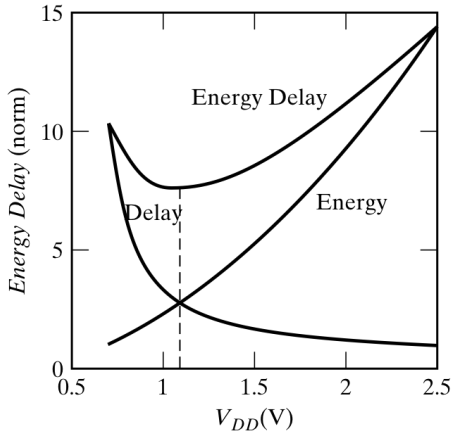

Intermezzo

Error-Correcting Codes

- Data transmissions, computations or storage may often fail
  - Because of yield issues
  - Noise
  - Interference
- Often used in wireless transmissions, hard disks, memory
- Redundancy can help to eliminate the impact of these errors
  - Triple-modular redundancy
  - Coding – add extra bits so that errors can be detected and corrected

Richard Wesley Hamming

- In telecommunication, a Hamming code is a linear error-correcting code named after its inventor, Richard Hamming. Hamming codes can detect up to two simultaneous bit errors, and correct single-bit errors; thus, reliable communication is possible when the Hamming distance between the transmitted and received bit patterns is less than or equal to one. By contrast, the simple parity code cannot correct errors, and can only detect an odd number of errors.
- Because of the simplicity of Hamming codes, they are widely used in computer memory (RAM). In particular, a single-error-correcting and double-error-detecting variant commonly referred to as SECDED.

R. Hamming, 1915-1998

Error-Correcting Codes - Example

Hamming Codes

$P_1 P_2 B_3 P_4 B_5 B_6 B_7$ *e.g. B3 Wrong*

with

$P_1 \oplus B_3 \oplus B_5 \oplus B_7 = 0$ $2^k \geq n+k+1$ 1

$P_2 \oplus B_3 \oplus B_6 \oplus B_7 = 0$ 1 *= 3*

$P_4 \oplus B_5 \oplus B_6 \oplus B_7 = 0$ 0

To correct 1 error in 4 bit word: 3 parity bits

In general: $2^k \geq k + n + 1$ (*n*: original bits; *k*: parity bits)
In this example: $8 \geq 3 + 4 + 1$

WE CAN DO BETTER

EE141 Class Project: LDPC Decoder

– from Algorithm to Circuit Implementation

Pulkit Grover, Stanley Chen, Nam-Seog Kim,
and Prof. Jan Rabaey

Spring, 2010

Motivation

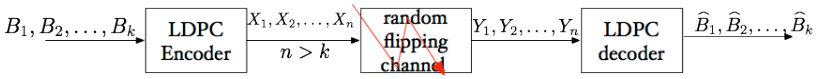

- Communication channel -- bit flipping error
- Introduce redundancy at the transmitter -- error correction coding
- Correct errors using redundancy at the decoder -- decoding
- A class of immense recent (theoretical and practical) interest : LDPC Codes

LDPC Codes

- A low-density parity-check (LDPC) code is a linear error correcting code, a method of transmitting a message over a noisy transmission channel. LDPC codes are capacity-approaching codes, which means that practical constructions exist that allow the noise threshold to be set very close (or even arbitrarily close on the BEC) to the theoretical maximum (the Shannon limit) for a symmetric memory-less channel. The noise threshold defines an upper bound for the channel noise up to which the probability of lost information can be made as small as desired. Using iterative belief propagation techniques, LDPC codes can be decoded in time linear to their block length.
- LDPC codes are finding increasing use in applications where reliable and highly efficient information transfer over bandwidth or return channel constrained links in the presence of data-corrupting noise is desired (10 GB Ethernet, DVBS-S, etc)
- LDPC codes are also known as Gallager codes, in honor of Robert G. Gallager, who developed the LDPC concept in his doctoral dissertation at MIT in 1960.

R. Gallager, 1931-

LDPC Iterative Decoding (5,3 code)

Check Node: Receive 5 bits from 5 Bit Nodes and send back parity check results

Major Operation: XOR, simple control logics

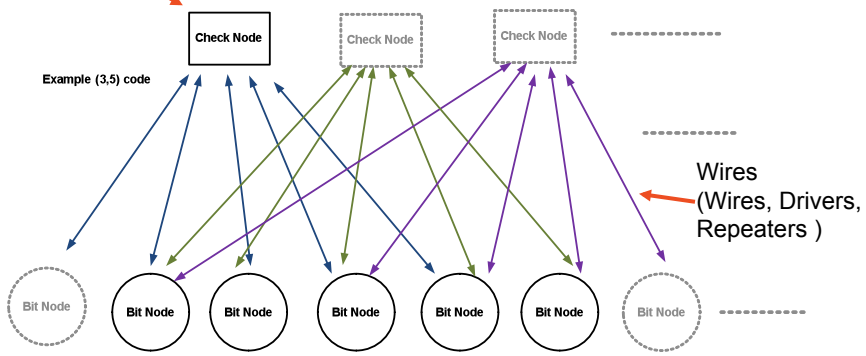

Bit Node: Send to 3 Check Nodes and receive the results from those Check Nodes

Major Operation: Majority voting, update value, simple control logics

LDPC Code

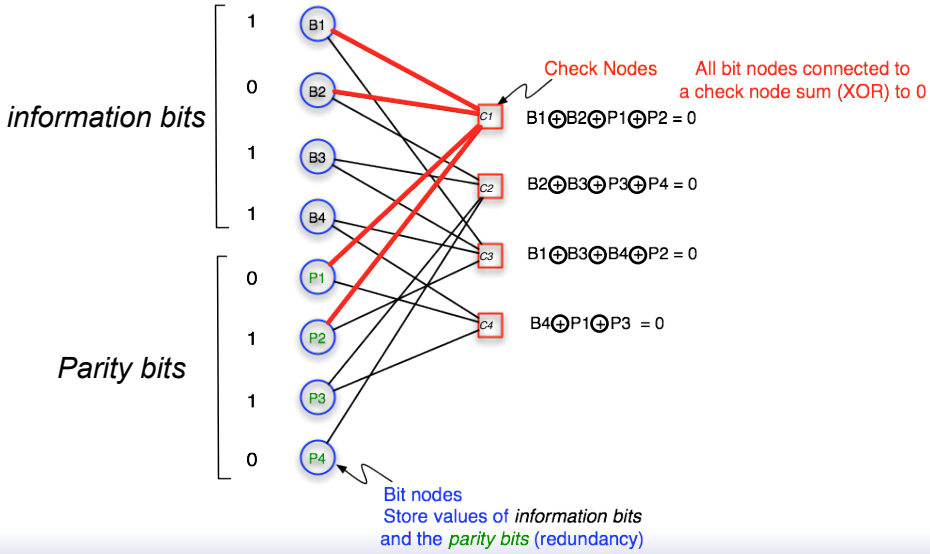

LDPC Decoding

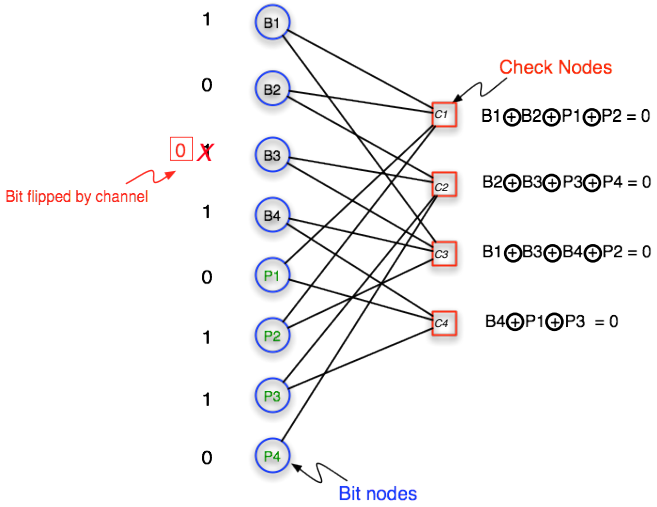

LDPC Decoding : First iteration

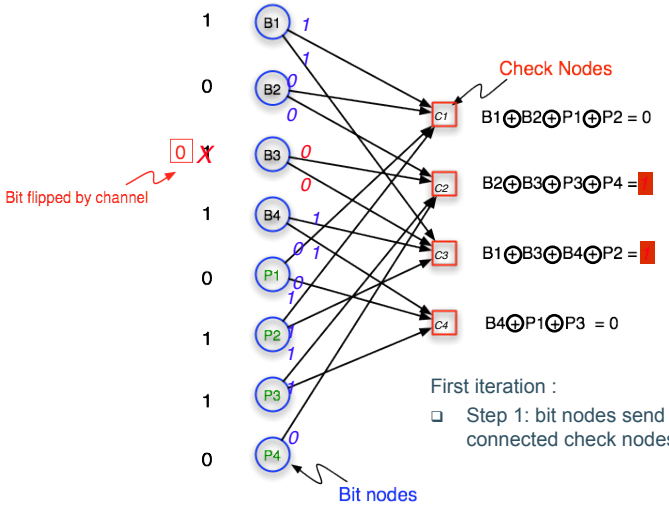

First iteration :

- Step 1: bit nodes send the channel outputs to the connected check nodes

LDPC Decoding : First iteration (contd.)

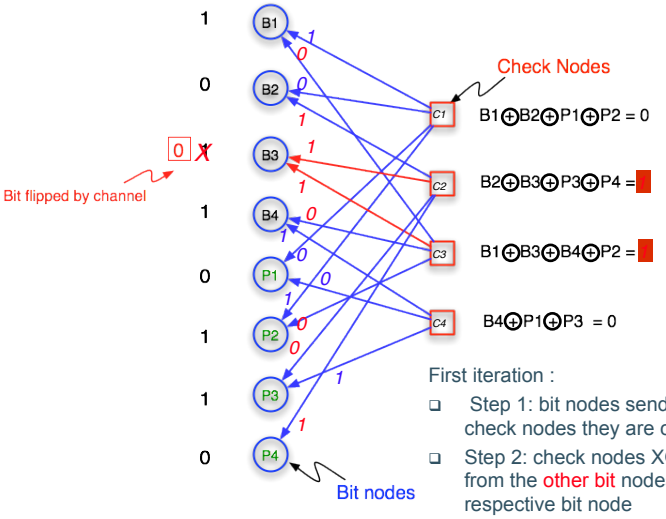

First iteration :

- Step 1: bit nodes send the channel outputs to check nodes they are connected to
- Step 2: check nodes XOR the values they received from the other bit nodes and send it to the respective bit node

LDPC Decoding : Second iteration

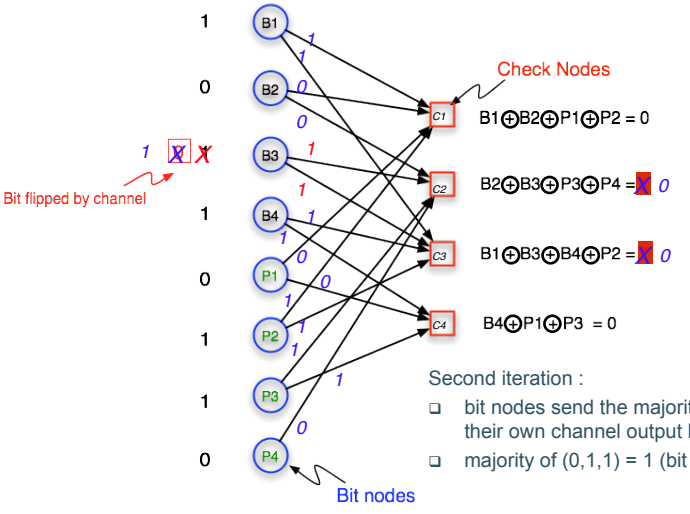

Second iteration :

- bit nodes send the majority of the messages and their own channel output back to the check nodes
- majority of (0,1,1) = 1 (bit decoded correctly)

Check Node

Check Node: Receive 4 (M) bits from Bit Nodes and send back parity check results

Major Operation: XOR, simple control logics

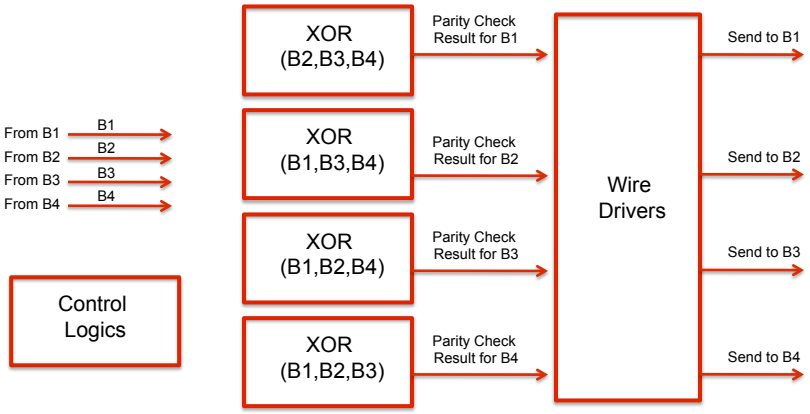

Bit Node

Bit Node: Receive from 2 (N) Check Nodes and channel and sent the result back to those Check Nodes

Major Operation: Majority voting, update value, simple control logics

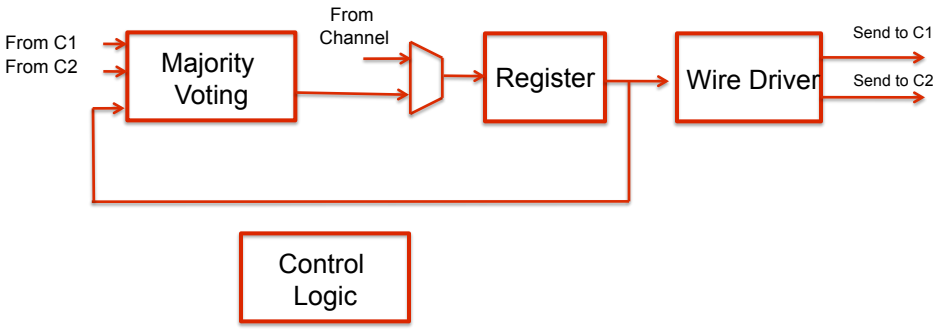

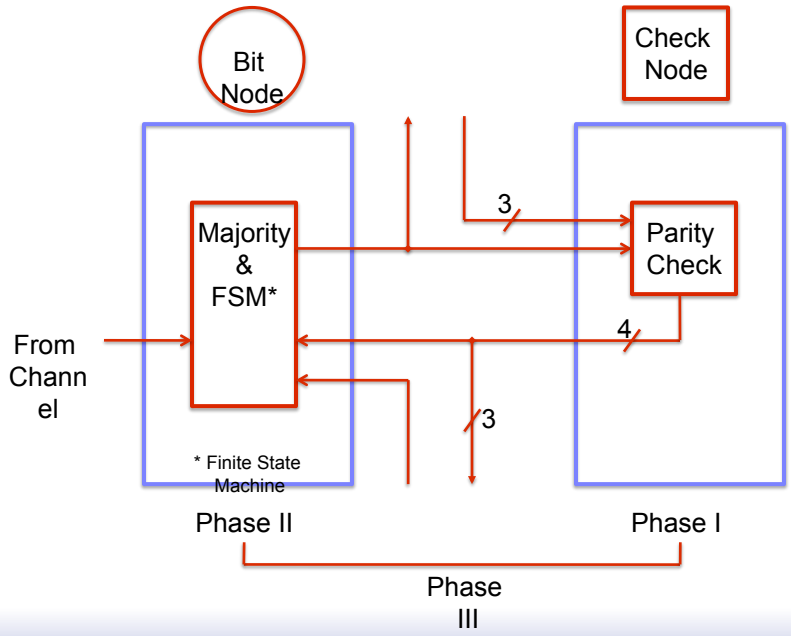

Project Overview

3-Phases

- Phase 1 (warm-up)
  - Introduction to LDPC decoding
  - Design a check node
- Phase 2 (deeper design experience)
  - Design of bit node
  - Complete design of a decoder (with help)
- Phase 3 (design optimization)
  - Optimize your design for power-performance tradeoffs
  - Go for the gold!

Time Line

Prj-ph1: 03/05~03/15, 11 days
Prj-ph2: 03/15~04/09, 25 days
Prj-ph3: 04/09~05/05, 26 days

6	11	Febr. 24	We	MOS cap + Inverter VTC	Lec11		Chapter 3,5
	12	Febr. 26	Fr	Mos cap + VTC revisited		Hw4	Chapter 3,5
7	13	March 3	We	CMOS Delay model, power			Chapter 5
	14	March 5	Fr	CMOS power		**Hw 5**	Chapter 5
8	15	March 10	We	Technology Scaling			Chapter 5,8
	16	March 12	Fr	Complex Logic			Chapter 6
9	17	March 17	We	Pass Transistor Logic			Chapter 6
	18	March 19	Fr	Dynamic Logic		**Hw 6**	Chapter 6
10	SPRING BREAK						
	19	March 31	We	Energy/Power			Chapter 6
	20	April 2	Fr	Memory		**Hw 7**	Chapter 12
11	21	April 7	We	MIDTERM II			Chapter 12
	22	April 9	Fr	Memory			Chapter 7
12	23	April 14	We	Latches/Registers			Chapter 7,10
	24	April 16	Fr	Sequential / Timing		**Hw 8**	Chapter 10
13	25	April 21	We	Timing			
	26	April 23	Fr	Clocks and Power		**Hw 9**	Chapter 10,9
14	27	April 28	We	Arithmetic			Chapter 11
	28	April 30	Fr	Perspectives		**Hw10 (optional)**	
15	May 3-7: RRR Week - Project Poster Presentations (BWRC)						
16		May 11	Tu	Final, 11:30am-2:30pm			

Prj-ph1 launched (Friday Mar. 5)
11 days
Prj-ph1 due (Monday Mar. 15) & Prj-ph2 launched
25 days
Prj-ph2 due &Prj-ph3 launched (Friday Apr. 9)
26 days
Poster Session (Wednesday May 5)

Project Phase 1 - Tasks

- Do an exercise on LDPC decoding
  - Given an input vector + decoding procedures, derive the final results step by step
- Derive the schematic for Check Node
- Design and size logic gates (w/ preset wire loading) for delay
- Simulate the design in Cadence SPECTRE

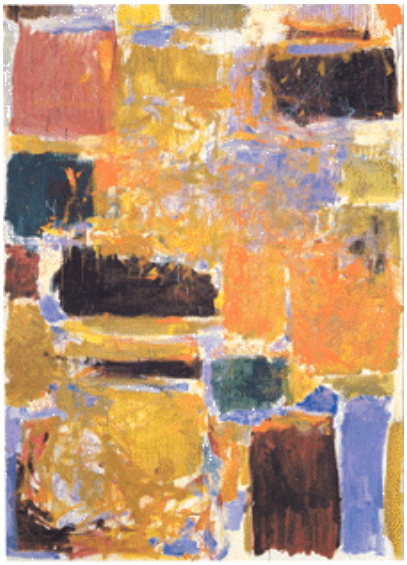

CMOS Logic Revisited

Analyzing and Optimizing Complex CMOS Gates

- Techniques very similar to the inverter case
- Logical Effort technique as the means for gate sizing and topology optimization
- However … some other things to be aware of

Fan-In Considerations

RC model:

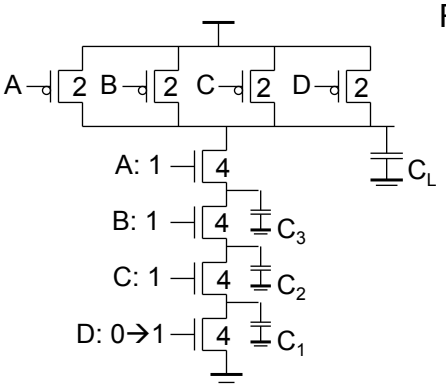

Fan-In Considerations

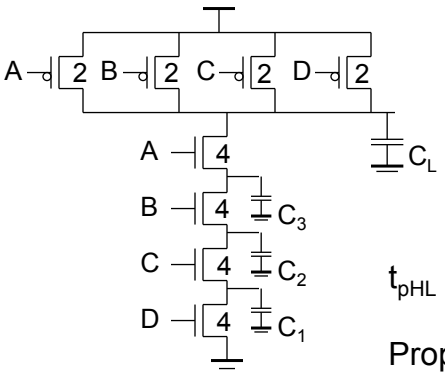

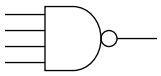

Distributed RC model (Elmore delay)

$t_{pHL} = 0.69\ R_{eqn}(C_1+2C_2+3C_3+4C_L)$

Propagation delay deteriorates rapidly as a function of fan-in – quadratically in the worst case.

t_p as a Function of Fan-In

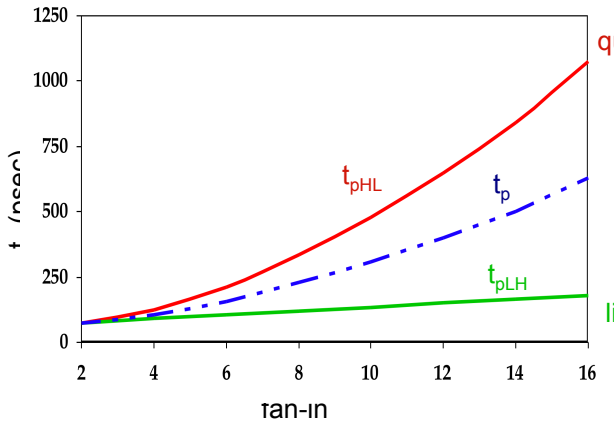

Gates with a fan-in greater than 4 should be avoided.

t_p as a Function of Fan-Out

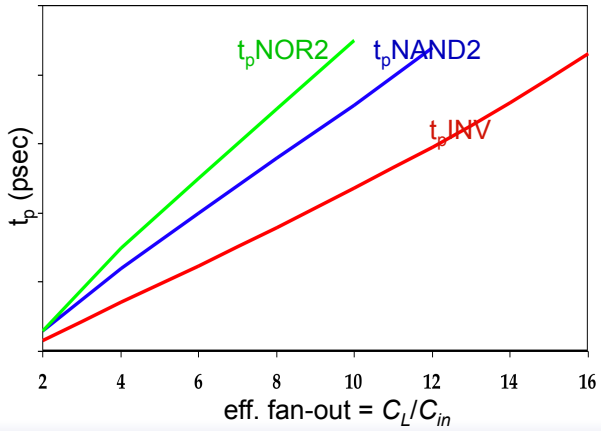

All gates have the same drive current.

Slope is a function of "driving strength"

t_p as a Function of Fan-In and Fan-Out

- Fan-in: quadratic due to increasing resistance and capacitance
- Fan-out: each additional fan-out gate adds two gate capacitances to C_L

$$t_p = a_1 FI + a_2 FI^2 + a_3 FO$$

Fast Complex Gates: Design Technique 1

- Transistor sizing
  - as long as fan-out capacitance dominates
- Progressive sizing

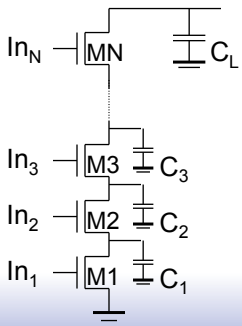

Distributed RC line

M1 > M2 > M3 > … > MN
(the FET closest to the output is the smallest)

Can reduce delay by more than 20%;
Be careful: input loading, junction caps, decreasing gains as technology shrinks

Fast Complex Gates: Design Technique 2

- Transistor ordering

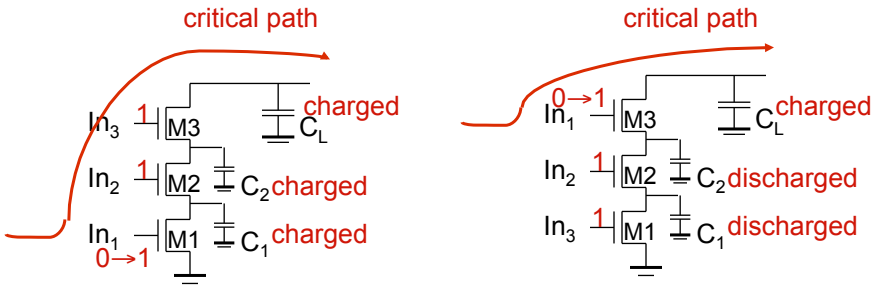

delay determined by time to discharge C_L, C_1 and C_2

delay determined by time to discharge C_L

Fast Complex Gates: Design Technique 3

- Alternate logic structures

F = ABCDEFGH

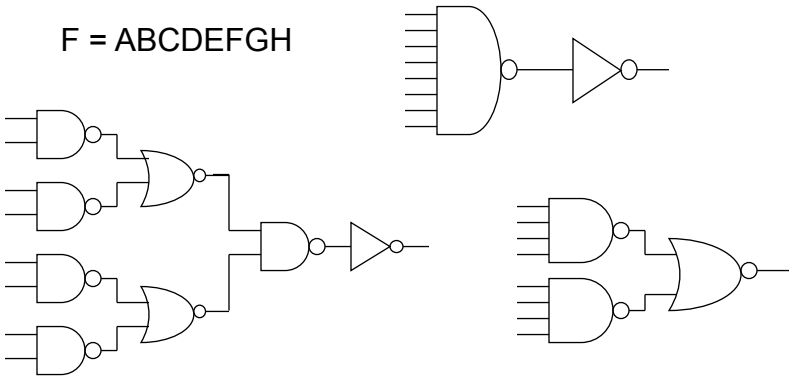

Fast Complex Gates: Design Technique 4

- Isolating fan-in from fan-out using buffer insertion

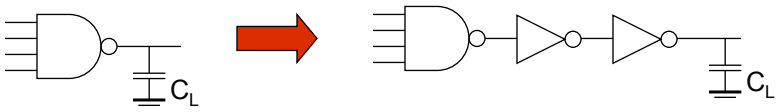

Fast Complex Gates: Design Technique 5

- Reducing the voltage swing

$$t_{pHL} = 0.5\,(C_L\,V_{DD}) / I_{DSATn}$$

$$= 0.5\,(C_L\,V_{swing}) / I_{DSATn}$$

  - linear reduction in delay
  - also reduces power consumption
- But the following gate is slower!
- Or requires use of "sense amplifiers" on the receiving end to restore the signal level (memory design)